

# Team Members:

**Pranjal Sahu** 110931955

**Santiago Vargas** 108721656

## Data Objects

*# User will create this object to send using a Client*

**UserRequest {**

**fields:**

objectId

operation

**}**

*# Client creates this object to send to a head replica*

**ClientRequest {**

**fields:**

objectId

operation

clientKey

requestId

**methods:**

constructor(obj\_id, op, key):

self.objectId = obj\_id

self.operation = op

self.clientKey = key

**}**

*# Result received from a tail replica*

**Result {**

**fields:**

requestID

result

list resultProof

**}**

**Order {**

**fields:**

operation

objectId

**}**

```

OrderProof {
  fields:
    proofs # list of proofs
    issuer # the issuer of this proof
}

```

```

Proof {
  fields:
    issuer          # the issuer of the proof which will be the replica id
    operation       # the operation which was performed
    objectId       # object on which operation was performed
    slot           # the unique operation number
}

```

## Classes

*# User will create a Client to send/receive multiple UserRequests*

```

Client {
  fields:
    key
    timeoutSpan
    list sent      # list of sent requests
    list received  # list of received responses
    requestCount

```

### **method declaration:**

```

    sendRequest(UserRequest)

    __sendRequest(UserRequest, msg, destReplicas)

    handleReplicaResponse(Result)

    handleError(Error)

```

### **implementation:**

```

constructor(tSpan):
    self.key = Olympus.obtainClientKey()
    self.timeoutSpan = tSpan

```

```

# User should use this interface
sendRequest(UserRequest):
    # Get a unique request id
    requestID = getUniqueReuquestID()

    # Get the current configuration
    config = Olympus.getCorrectConfig()

    # Send a request to the head replica
    dest = config.head
    return __sendRequest(UserRequest, msg = "newRequest", destReplicas = dest)

__sendRequest(UserRequest, msg, destReplicas):
    config = Olympus.getCorrectConfig()

    # Create a request
    request = new ClientRequest(obj_id = UserRequest.objectId, op =
UserRequest.operation, key = self.key)

    # Send the request to the destination replicas
    send sign(<msg, request>) to destReplicas    # Sign with the replica's public key
    sent.add(request)

    # The request has been sent and wait for at max timeout interval for the result
    timeout(minutes = self.timeoutSpan):
        # Check if we have received
        if request.requestID not in received:
            # Get the current config. This updates the config incase there are any changes
            config = Olympus.getCorrectConfig()
            self.__sendRequest(UserRequest, msg = 'retranmission', destReplicas =
config.getAllReplicas())

# Receive an unvalidated result
# Note: We assume a user may get multiple responses for a single request
handleReplicaResponse(Result):
    if Result.requestId in sent:
        for each proof in Result.resultProofs
            replicaPublicKey = config.replicas(proof.replicaID).publicKey
            if SHA256(Result.result, replicaPublicKey) != proof.signature
                return error

```

```
return Result.result
```

```
# Received an error. This means the sender replica is immutable
```

```
handleError(Error):
```

```
    config = Olympus.getCorrectConfig()
```

```
    __sendRequest(UserRequest, replicas = config.getAllReplicas())
```

```
}
```

```
Olympus {
```

```
    fields:
```

```
        config                # Ordered list of replicas
```

```
        replicaCount          # The number of replicas
```

```
        list wedge_replies     # list of wedge replies obtained from replicas
```

```
        list catchUpReplies    # list of catch up replies obtained from replicas
```

```
        list runningStateReplies # list of running state replies from replicas
```

```
        list publicReplicaKeys # list of public keys for replica
```

```
        list privateReplicaKeys # list of private keys for replicas
```

```
        replyTimeout           # acceptable amount of time to wait for a reply from a replica
```

```
    methods:
```

```
        initHist(numberOfReplicas, config, hist, runningState)
```

```
        receiveClientKeyRequest()
```

```
        getCorrectConfig()
```

```
        checkValidWedgeReply(wedge_reply)
```

```
        validateOrderProof(orderProofs)
```

```
        handleAsyncRequest(msg, request, client)
```

```
        handleWedgeResponse(wedgeResponse)
```

```
        handleCatchupResponse(catchupResponse)
```

```
        handleRunningStateResponse(runningStateResponse)
```

```
        validateCheckpointProof(checkpointProof)
```

```
        receiveReconfigureRequest()
```

```

constructor(numberOfReplicas, timoeut):
    self.replicaCount = numberOfReplicas
    self.config = new Config()
    # Initialize a configuration with empty history and empty running state
    initHist(self.replicaCount, self.config, hist = {}, runningState = {})
    # Generate new public and private keys
    self.publicReplicaKeys, self.privateReplicaKeys = generateSignKeys(replicaCount)
    self.replyTimeout = timeout

```

*# Creates a new configuration*

```

initHist(numberOfReplicas, config, hist, runningState):
    # Delete old replicas
    config.purgeReplicas()
    # Generate new public and private keys
    self.publicReplicaKeys, self.privateReplicaKeys = generateSignKeys(replicaCount)

```

*# Create each replica*

```

for id in numberOfReplicas:
    r = new Replica(publicKeys = self.publicReplicaKeys, secretSigningKey =
                    self.privateReplicaKeys[id], status = Active, history = hist,
                    runningState = runningState)
    # Add the replica to the config
    config.add(r)

```

*# Link each replica*

```

for each replica in config:
    if replica is not the head:
        replica.previousReplica = replica.index-1
        replica.headReplica = first replica in config
    if replica is not the tail:
        replica.nextReplica = replica.index+1
    replica

```

*# Returns a new identifier and key to the client*

```

receiveClientKeyRequest():
    return generateRandomKey()

```

*# Asynchronous method to handle requests from Clients to obtain the replica list*

```

getCorrectConfig():
    return self.config

```

```

checkValidWedgeReply(wedge_reply):
    for order_proof in wedge_reply.order_proofs:
        replicald = order_proof.issuerId
        signKey = signKeys[replicald][replicald+1]
        if (unsign(signKey.privateKey, order_proof.data) == Success )
            order_proof_data = unsign(signKey.privateKey, order_proof.data)
        else:
            return False

validateOrderProof(orderProofs):
    for order_proof in order_proofs:
        # Check source of order_proof by using replica's public key
        if decrypt(order_proof) == False:
            return False
        # Check if only one distinct operation exists in this order proof
        if order_proof.operations.distinct.count != 1
            return False # Not correct, therefore false
        if order_proof.slots.distinct.count != 1 # Check if only one slot is present
            return False
    return True

# Asynchronous handler for all messages. This multiplexes to sub-handlers
handleAsyncRequest(msg, request, client):
    switch msg:
        'wedgeResponse':
            handleWedgeResponse(request)
        'catchupResponse':
            handleCatchupResponse(request)
        'runningStateResponse':
            handleRunningStateResponse(request):
        'reconfigure':
            receiveReconfigureRequest()
        'configuration':
            getCorrectConfig()
    default:
        return Error

```

*# Handler for wedge responses from replicas*

**handleWedgeResponse**(wedgeResponse):

*# Append all the wedge responses obtained*

wedge\_replies.append(wedgeResponse)

*# Handler for catch up messages from replicas*

**handleCatchupResponse**(catchupResponse):

*# Append all the catch up responses obtained*

catchUpReplies.append(catchupResponse)

*# Handler for runningState messages from replicas*

**handleRunningStateResponse**(runningStateResponse):

*# Append all the running state responses obtained*

runningStateReplies.append(runningStateResponse)

*# All the hashes should match in a checkpointProof*

**validateCheckpointProof**(checkpointProof):

hash h\_hash = {}

hash h\_seq = {}

for proof in checkpointProof

if decrypt(proof) == True:

h\_hash[proof.hash] = 1

*# There is only hash present in all the proofs inside checkpointproof*

if h\_hash.keys().size > 1

return False

*# Sequence number of all the proof in checkPointProof should be same*

if h\_seq.keys().size > 1

return False

return True

**receiveReconfigureRequest**():

*# Request wedge statements from each replica*

wedge\_replies = []

for each replica in config:

*# Signing uses the private key of the Olympus*

send sign(<"wedgeRequest">) to replica

*# Wait untill obtaining t+1 wedge replies from the replica*

*# This will create a quorum of t+1 replicas which will ensure that there is at least one*

*# honest replica*

While (wedge\_replies.size < (replicaCount - 1)/2 + 1):

    True *# Keep looping*

*# Find the maximal orderproofs*

*# slot\_op contains the maximum length of order proof for each slot and the operation  
    done in each slot*

*# Create a list of [slot, operation] tuple to recreate replica*

hash slot\_op = {} *# Length and operation for each slot*

for wedge\_reply in wedge\_replies:

*# Check valid source of wedge\_reply using the public key for that replica*

    if (decrypt(wedge\_reply) == True)

*# Check valid order\_proofs and checkpoint proofs*

        if(validateOrderProof(wedge\_reply.order\_proofs)

            and validateCheckpointProof(wedge\_reply.checkpoint\_proof)):

            if slot\_op[wedge\_reply.order\_proof.slot] is empty:

                slot\_op[slot] = [wedge\_reply.order\_proof.operation,  
                                    wedge\_reply.order\_proof.length]

            else if slot\_op[slot] is present:

                if slot\_op[slot].length < wedge\_reply.order\_proof.length

                    slot\_op[slot] = [wedge\_reply.order\_proof.operation,  
                                    wedge\_reply.order\_proof.length]

            else

                continue

        else

            continue

    else:

        continue

*# Creating history to be used for reconfiguring the replicas*

hash h = {}

for slot in slot\_op:

    h[slot] = slot\_op[slot].operation

for wedge\_reply in wedge\_replies:

    validateCheckpointProof(wedge\_reply.checkpoint\_proof)



```

# Create catch_up messages content for checkpoint purpose
for replica in config: # For each replica
    # Calculate the difference between validated slot operations and the replica's history
    (wedge_reply.order_proofs)
    catchupDifference = wedge_replies[replica.id].order_proofs.keys - slot_op.keys
    # Signing uses the private key of the Olympus
    send sign(<"catchUpRequest", catchupDifference>) to replica

# Wait an acceptable amount of time for catchup responses
# If a replica does not reply in an acceptable amount of time then they are faulty
# Therefore, we are guaranteed at least t + 1 responses since t + 1 replicas are correct
catchUpReplies.purge() # Clear the reply list
timeout(self.replyTimeout):
    # Dictionary of catchup responses (running state hashes) with a default value of 0
    hash stateHashes = {}

# Once the timeout is over, go through the list of response hashes
for catchUpResponse in self.catchUpReplies:
    # Decrypt the hash using the replica's public key. This validates the replica's identity
    decrypt(catchUpResponse)
    stateHashes[catchUpResponse.stateHash] += 1

stateHashes.sort() # Sort by value

# Since the list is sorted, the last hash should have the largest and should have a value of
# at least t + 1 (at least t+1 replicas have the same hash)
correctHash = stateHash.last().key
for catchUpResponse in catchUpReplies: # Already decrypted
    if catchUpResponse.stateHash == correctHash:
        self.runningStateReplies.purge() # Clear the response list
        send sign(<"runningStateRequest">) to catchUpResponse.replica
        # If we do not receive a response then try another replica
        timeout(self.replyTimeout):
            for stateResponse in self.runningStateReplies:
                # Decrypt using the replica's public key and check if the hashes match
                if decrypt(stateResponse) and if SHA256(stateResponse.runningState) ==
correctHash:
                    correctRunningState = stateResponse.runningState
                    break(2) # Break 2 levels - out of the outer loop

```

```

        # Create a new configuration
        initHist(self.replicaCount, self.config, hist = slot_op, runningState = correctRunningState)
    }

```

**Replica {**

***fields:***

```

    previousReplica
    nextReplica
    headReplica          # the address to head Replica
    history
    interReplicaKeys     # public keys of all replicas
    privateKey           # the private key of this replica
    status
    cache                # cache for storing the results
    headFlag             # boolean flag to check if replica is head or not
    tailFlag             # boolean flag to check if replica is tail or not
    timeoutSpan          # the timeout for retransmission requests
    pendingRequests      # requests whose result shuttle is pending
    allowedOperations = ['query','update' ]
    maxSlotNum           # the largest slot number ordered
    runningState         # contains the values of all objects in the replica
    checkpointNum        # the latest slot number which has been checkpointed
    checkpointHash       # the hash of the running state at the latest checkpoint
    checkPointProof      # latest checkpoint Proof

```

***methods:***

```

    checkpoint(checkPointSlot)

    handleAsyncRequest(msg, request, client)

    handleNewRequest(client, request)

    handleRetranmissionRequest(client, request)

    handleWedgeRequest(request)

    handleResultShuttle(result, resultProof)

    handleCheckpointShuttle(CheckPointProof)

    handleCheckpointReturnShuttle(CheckPointProof)

```

```
handleCatchupRequest(request)
```

```
handleRunningStateRequest(request)
```

```
getProofsFromHistory(required_slot)
```

```
orderCommand(order, orderProof, resultProof, sendFlag = True)
```

```
constructor(publicReplicaKeys, secretSigningKey, status, history, runningState):
```

```
    self.status = status
```

```
    self.interReplicaKey = publicReplicaKeys
```

```
    self.privateKey = secretSigningKey
```

```
    self.maxSlotNum = max(history) # Find the maximum slot number in the history
```

```
    self.checkpointInterval = 100 # Denotes how often to checkpoint in terms of  
                                # new slots ie. every 100 slots
```

```
    self.runningState = runningState
```

```
# This function performs a checkpoint up to the checkPointSlot
```

```
checkpoint(checkPointSlot):
```

```
    # Update checkpointSlot number
```

```
    self.checkpointNum = checkPointSlot
```

```
    # Hash the running state
```

```
    self.checkPointHash = SHA256(self.runningState) # Calculate hash using SHA256
```

```
# Catch all message handler
```

```
handleAsyncRequest(msg, request, client):
```

```
    switch msg:
```

```
        'newRequest':
```

```
            handleNewRequest(client, request)
```

```
        'retransmission':
```

```
            handleRetranmissionRequest(client, request)
```

```
        'wedgeRequest':
```

```
            handleWedgeRequest(request)
```

```
        'catchUpRequest':
```

```
            handleCatchupRequest(request)
```

```
        'runningStateRequest':
```

```
            handleRunningStateRequest(request)
```

```
        'returnShuttle':
```

```

        handleResultShuttle(result, resultProof)
    'checkpointShuttle':
        handleCheckpointShuttle(CheckPointProof)
    'checkpointReturnShuttle':
        handleCheckpointReturnShuttle(CheckPointProof)
    default:
        send sign(<"error", request.requestid>) to client      # Sign using replica's private key

# Asynchronous handler for normal order command requests
# Multi purpose handler to handle new requests
handleNewRequest(client, request):

    if self.state != ACTIVE:                                     # Check if we are active
        send sign(<"error", request.requestid>) to client      # Sign using replica's private key
        return

    if valid(client.clientKey) == False:                         # Check validity of client key
        send sign(<"error", request.requestid>) to client      # Sign using replica's private key
    return
    order = new Order(request.operation, request.objectId)
    orderCommand(order)

# Handler for retransmission requests
handleRetranmissionRequest(client, request):
    if self.status == Immutable:
        send sign(<"error", request.requestid>) to client      # Sign using replica's private key
        return
    else if request.id in cache:                                  # Sign using replica's private key
        send sign(<"retranmissionResponse", cache[request.id]>) to client
        return
    else:
        if self.headFlag == False:                               # Forward the request to head
            sendRequestToReplica(headReplica, request)
            timeout(minutes = self.timeoutSpan):
                if request.id not in cache:                       # If no result shuttle obtained, then reconfig
                    self.status = Immutable
                    send sign(<"reconfigure">) to Olympus        # Sign using replica's private key
                    return
            Else
                # Sign using replica's private key
                send sign(<"retranmissionResponse", cache[request.id]>) to client

```

```

        return
    else:
        # We are the head replica
        if request.id in cache:
            # If no result shuttle obtained, then reconfig
            # Sign using replica's private key
            send sign(<"retranmissionResponse", cache[request.id]>) to client
            return
        else if request.id in pendingRequests:
            timeout(minutes = self.timeoutSpan):
                if request.id not in cache
                    # If no result shuttle obtained, then reconfig
                    self.status = Immutable
                    send sign(<"reconfigure">) to Olympus # Sign using replica's private key
                    return
                Else
                    # Sign using replica's private key
                    send sign(<"retranmissionResponse", cache[request.id]>) to client
                    return
        else:
            # Similar to new request, therefore calling handleNewRequest
            handleNewRequest(client, request)

# Handler for wedge requests
handleWedgeRequest(request):
    if (decrypt(request) == False)
        # Using Olympus public key
        return False
        # Do not respond

    # Make sure our status is Immutable
    self.status = Immutable

    # Sign the data and send it to Olympus
    send <"wedgeResponse", self.runningState, self.orderProofs, self.checkPointProof>
    to Olympus

# Handler for a result shuttle
handleResultShuttle(result, resultProof):
    validFlag = True
    for proof in resultProof:
        if valid(proof) != True
            validFlag = False
        break
    if validFlag == False:
        # If the return Shuttle is not valid, then ignore
        return
    else:

```

```
cache[result.id] = result           # If it is valid, then cache the result  
pendingRequests.delete(result.id)  # Delete from the list of pending requests
```

*# Handler to forward the checkpoint shuttle to the next replica*

```
handleCheckpointShuttle(CheckPointProof):  
    checkPointTuple = <self.checkpointNum, self.checkPointHash>  
    CheckPointProof.append(checkpointTuple)  
    if self.tailFlag:  
        # Return to the previous replica  
        handleCheckpointReturnShuttle(CheckPointProof)  
    else:  
        # Forward to the next replica  
        send <"checkpoint", CheckPointProof> to self.nextReplica
```

*# Handler to validate the return checkpoint shuttle*

```
handleCheckpointReturnShuttle(CheckPointProof):  
    # Validation  
    for checkpointTuple in CheckPointProof:  
        if (checkpointTuple.checkpointNum != self.checkpointNum) or  
            (checkpointTuple.checkPointHash != self.checkPointHash):  
            return Error           # Do nothing  
  
    # Delete checkpointed proofs from history  
    checkpointSlot = CheckPointProof.checkpointTuple.checkpointNum  
    delete history[:checkpointSlot]    # Delete history till checkpoint  
  
    # Store the returned checkPointProof from the tail  
    self.checkPointProof = CheckPointProof  
  
    # Return to the previous replica if you are not the head  
    if not self.headFlag:  
        send <CheckPointProof> to self.previousReplica
```

*# Handler for catch up messages from Olympus*

```
handleCatchupRequest(request):  
    # Validate that this request came from Olympus. Decrypt using Olympus' public key  
    if not decrypt(request)  
        return False  
    for order in request.catchupDifference:
```

```

    # Order command and do not forward to the next replica
    orderCommand(order, new OrderProof, new ResultProof, sendFlag = False)

stateHash = SHA-256(self.runningState)
send sign(<"catchupResponse", stateHash>) to Olympus # Sign using replica's private key

# Handler for running state requests from Olympus
handleRunningStateRequest(request):
    # Validate that this request came from Olympus. Decrypt using Olympus' public key
    if not decrypt(request)
        return False
    send sign(<"runningStateResponse", self.runningState>) to Olympus

# Gets all relevant proofs from the history
# Start from the current checkpoint and reach the required slot, using operations in history
getProofsFromHistory(required_slot):
    proof_statements = self.history[sel.checkpointNum:required_slot]
    return proof_statements

# Relates to the orderCommand transition at the replica.
# Assumes we are active -> This is checked in the new request handler
orderCommand(order, orderProof = [], resultProof = [], sendFlag = True):
    # Validate history
    if self.history has slot: # A slot for this command already exists. Cannot issue it twice
        return False

    # Check if performing an allowed operation
    If order.operation not in self.allowedOperations:
        return False

    # Generate a new slot if we are the head
    if self.headFlag == True:
        sNum = self.maxSlotNum+1
    else
        sNum = order.slotNumber

    # Prevent holes in slot numbers
    if self.history does not contain sNum-1:checkpointNum: # We found a hole
        return False

```

```

# Update maxSlotNum
self.maxSlotNum = sNum

# Apply running state
result = runningState[request.objectId].apply(request.operation)

# Create proofs
order = new OrderProof(slotNumber = sNum,
                        operation = request.operation, proofs =
                                getProofsFromHistory(sNum))
signedOrder = sign(key = self.privateKey, statement = order)
orderProof.append(signedOrder)

signedResult = sign(key = self.privateKey, statement = result)
resultProof.append(signedResult)

# Initiate checkpoints after every X orders where X = self.checkpointInterval
if maxSlotNum is a multiple of self.checkpointInterval:
    checkpoint(maxSlotNum)
    if self.headFlag # If we are the head replica, initiate the checkpoint shuttle
        handleCheckpointShuttle(new CheckPointProof())

if not replica.tailFlag:
    send <"order", order, orderProof, resultProof> to self.nextReplica
Else:
    # send result to client
    send <"result", requestID, result, resultProof> to request.client
    # send the return shuttle
    send <"ReturnShuttle", result, resultProof> to self.previousReplica
}

```