

# Team Members:

**Pranjal Sahu**     110931955  
**Santiago Vargas**

## Data Objects

*#User will create this object to send using a Client*

```
UserRequest {  
  fields:  
    objectId  
    operation  
}
```

*#Client creates this object to send to a head replica*

```
ClientRequest {  
  fields:  
    objectId  
    operation  
    clientKey  
    requestId  
  
  methods:  
    constructor(obj_id, op, key):  
      self.objectId = obj_id  
      self.operation = op  
      self.clientKey = key  
}
```

*#Result received from a tail replica*

```
Result {  
  fields:  
    requestID  
    result  
    list resultProof  
}
```

```

OrderProof {
  fields:
    proofs # list of proofs
    issuer # the issuer of this proof
}

```

```

Proof {
  fields:
    issuer          # the issuer of the proof which will be the replica id
    operation       # the operation which was performed
    objectId        # object on which operation was performed
    slot            # the unique operation number
}

```

## Classes

*#User will create a Client to send/receive multiple UserRequests*

```

Client {
  fields:
    key
    timeoutSpan
    list sent      # list of sent requests
    list received  # list of received responses
    requestCount

```

### ***method declaration:***

```

    sendRequest(UserRequest)

    __sendRequest(UserRequest, msg, destReplicas)

    handleReplicaResponse(Result)

    handleError(Error)

```

### ***implementation:***

```

  constructor(tSpan):
    self.key = Olympus.obtainClientKey()
    self.timeoutSpan = tSpan

```

```

#User should use this interface
sendRequest(UserRequest):
    # Get a unique request id
    requestID = getUniqueReuquestID()

    # Get the current configuration
    config = Olympus.getCorrectConfig()

    # Send a request to the head replica
    dest = config.head
    return __sendRequest(UserRequest, msg = "newRequest", destReplicas = dest)

__sendRequest(UserRequest, msg, destReplicas):
    config = Olympus.getCorrectConfig()

    # Create a request
    request = new ClientRequest(obj_id = UserRequest.objectId, op =
UserRequest.operation, key = self.key)

    # Send the request to the destination replicas
    send sign(<msg, request>) to destReplicas    # Sign with the replica's public key
    sent.add(request)

    # The request has been sent and wait for at max timeout interval for the result
    timeout(minutes = self.timeoutSpan):
        # Check if we have received
        if request.requestID not in received:
            # Get the current config. This updates the config incase there are any changes
            config = Olympus.getCorrectConfig()
            self.__sendRequest(UserRequest, msg = 'retranmission', destReplicas =
config.getAllReplicas())

# Receive an unvalidated result
# Note: We assume a user may get mulitple responses for a single request
handleReplicaResponse(Result):
    if Result.requestId in sent:
        for each proof in Result.resultProofs
            replicaPublicKey = config.replicas(proof.replicaID).publicKey
            if SHA256(Result.result, replicaPublicKey) != proof.signature
                return error

```

```
return Result.result
```

```
# Received an error. This means the sender replica is immutable
```

```
handleError(Error):
```

```
    config = Olympus.getCorrectConfig()
```

```
    __sendRequest(UserRequest, replicas = config.getAllReplicas())
```

```
}
```

```
Olympus {
```

```
    fields:
```

```
        config                # Ordered list of replicas
```

```
        replicaCount          # The number of replicas
```

```
        list wedge_replies     # list of wedge replies obtained from replicas
```

```
        list catchUpReplies    # list of catch up replies obtained from replicas
```

```
        list runningStateReplies # list of running state replies from replicas
```

```
        list publicReplicaKeys # list of public keys for replica
```

```
        list privateReplicaKeys # list of private keys for replicas
```

```
        replyTimeout           # acceptable amount of time to wait for a reply from a replica
```

```
    methods:
```

```
        initHist(numberOfReplicas, config, hist, runningState)
```

```
        receiveClientKeyRequest()
```

```
        getCorrectConfig()
```

```
        checkValidWedgeReply(wedge_reply)
```

```
        validateOrderProof(orderProofs)
```

```
        handleAsyncRequest(msg, request, client)
```

```
        handleWedgeResponse(wedgeResponse)
```

```
        handleCatchupResponse(catchupResponse)
```

```
        handleRunningStateResponse(runningStateResponse)
```

```
        validateCheckpointProof(checkpointProof)
```

```
        receiveReconfigureRequest()
```

```

constructor(numberOfReplicas, timoeut):
    self.replicaCount = numberOfReplicas
    self.config = new Config()
    # Initialize a configuration with empty history and empty running state
    initHist(self.replicaCount, self.config, hist = {}, runningState = {})
    # Generate new public and private keys
    self.publicReplicaKeys, self.privateReplicaKeys = generateSignKeys(replicaCount)
    self.replyTimeout = timeout

# Creates a new configuration
initHist(numberOfReplicas, config, hist, runningState):
    # Delete old replicas
    config.purgeReplicas()
    # Generate new public and private keys
    self.publicReplicaKeys, self.privateReplicaKeys = generateSignKeys(replicaCount)

    # Create each replica
    for id in numberOfReplicas:
        r = new Replica(publicKeys = self.publicReplicaKeys, secretSigningKey =
self.privateReplicaKeys[id], status = Active, history = hist, runningState = runningState)
        # Add the replica to the config
        config.add(r)

    # Link each replica
    for each replica in config:
        if replica is not the head:
            replica.previousReplica = replica.index-1
            replica.headReplica = first replica in config
        if replica is not the tail:
            replica.nextReplica = replica.index+1
    replica

# Returns a new identifier and key to the client
receiveClientKeyRequest():
    return generateRandomKey()

# Asynchronous method to handle requests from Clients to obtain the replica list
getCorrectConfig():
    return self.config

```

```

checkValidWedgeReply(wedge_reply):
    for order_proof in wedge_reply.order_proofs:
        replicald    = order_proof.issuerId
        signKey      = signKeys[replicald][replicald+1]
        if (unsign(signKey.privateKey, order_proof.data) == Success )
            order_proof_data = unsign(signKey.privateKey, order_proof.data)
        else:
            return False

validateOrderProof(orderProofs):
    for order_proof in order_proofs:
        # Check source of order_proof by using replica's public key
        if decrypt(order_proof) == False:
            return False
        # Check if only one distinct operation exists in this order proof
        if order_proof.operations.distinct.count != 1
            return False                                # Not correct, therefore false
        if order_proof.slots.distinct.count != 1          # Check if only one slot is present
            return False
    return True

# Asynchronous handler for all messages. This multiplexes to sub-handlers
handleAsyncRequest(msg, request, client):
    switch msg:
        'wedgeResponse':
            handleWedgeResponse(request)
        'catchupResponse':
            handleCatchupResponse(request)
        'runningStateResponse':
            handleRunningStateResponse(request):
        'reconfigure':
            receiveReconfigureRequest()
        'configuration':
            getCorrectConfig()
    default:
        return Error

# Handler for wedge responses from replicas
handleWedgeResponse(wedgeResponse):
    # Append all the wedge responses obtained

```

```
wedge_replies.append(wedgeResponse)
```

```
# Handler for catch up messages from replicas
```

```
handleCatchupResponse(catchupResponse):
```

```
    # Append all the catch up responses obtained  
    catchUpReplies.append(catchupResponse)
```

```
# Handler for runningState messages from replicas
```

```
handleRunningStateResponse(runningStateResponse):
```

```
    # Append all the running state responses obtained  
    runningStateReplies.append(runningStateResponse)
```

```
# All the hashes should match in a checkpointProof
```

```
validateCheckpointProof(checkpointProof):
```

```
    hash h_hash = {}
```

```
    hash h_seq = {}
```

```
    for proof in checkpointProof
```

```
        if decrypt(proof) == True:
```

```
            h_hash[proof.hash] = 1
```

```
    # There is only hash present in all the proofs inside checkpointproof
```

```
    if h_hash.keys().size > 1
```

```
        return False
```

```
    # Sequence number of all the proof in checkPointProof should be same
```

```
    if h_seq.keys().size > 1
```

```
        return False
```

```
    return True
```

```
receiveReconfigureRequest():
```

```
    # Request wedge statements from each replica
```

```
    wedge_replies = []
```

```
    for each replica in config:
```

```
        # signing uses the private key of the Olympus
```

```
        send sign(<"wedgeRequest">) to replica
```

```
    # Wait till obtaining t+1 wedge replies from the replica
```

```
    # This will create a quorum of t+1 replicas which will ensure that there is at least one  
    honest replica
```

```
    While (wedge_replies.size < (replicaCount - 1)/2 + 1):
```

```
        True
```

```
        # Keep looping
```

```

# Find the maximal orderproofs
# slot_op contains the maximum length of order proof for each slot and the operation
  done in each slot
# Create a list of [slot, operation] tuple to recreate replica

```

```

hash slot_op = {}          # Length and operation for each slot
for wedge_reply in wedge_replies:
    # Check valid source of wedge_reply using the public key for that replica
    if (decrypt(wedge_reply) == True)
        # Check valid order_proofs and checkpoint proofs
        if(validateOrderProof(wedge_reply.order_proofs) and
validateCheckpointProof(wedge_reply.checkpoint_proof)):

            if slot_op[wedge_reply.order_proof.slot] is empty:
                slot_op[slot] = [wedge_reply.order_proof.operation,
wedge_reply.order_proof.length]

            else if slot_op[slot] is present:

                if slot_op[slot].length < wedge_reply.order_proof.length
                    slot_op[slot] = [wedge_reply.order_proof.operation,
wedge_reply.order_proof.length]
                else
                    continue
            else
                continue
        else:
            continue

hash h = {}                # Creating history to be used for reconfiguring the replicas
for slot in slot_op:
    h[slot] = slot_op[slot].operation

for wedge_reply in wedge_replies:
    validateCheckpointProof(wedge_reply.checkpoint_proof)

# Create catch_up messages content for checkpoint purpose
for replica in config: # For each replica
    # Calculate the difference between validated slot operations and the replica's history
    (wedge_reply.order_proofs)

```



```

        catchupDifference = wedge_replies[replica.id].order_proofs.keys - slot_op.keys
        send sign(<"catchUpRequest", catchupDifference>) to replica          # signing uses
the private key of the Olympus

```

```

# Wait an acceptable amount of time for catchup responses

```

```

# If a replica does not reply in an acceptable amount of time then they are faulty

```

```

# Therefore, we are guaranteed at least t + 1 responses since t + 1 replicas are correct

```

```

catchUpReplies.purge()          # Clear the reply list

```

```

timeout(self.replyTimeout):

```

```

    # Dictionary of catchup responses (running state hashes) with a default value of 0

```

```

    hash stateHashes = {}

```

```

    # Once the timeout is over, go through the list of response hashes

```

```

    for catchUpResponse in self.catchUpReplies:

```

```

        decrypt(catchUpResponse)      # Decrypt the hash using the replica's public key.

```

```

This validates the replica's identity

```

```

        stateHashes[catchUpResponse.stateHash] += 1

```

```

stateHashes.sort()              # Sort by value

```

```

correctHash = stateHash.last().key    # Since the list is sorted, the last hash should
have the largest and should have a value of at least t + 1 (at least t+1 replicas have the same
hash)

```

```

for catchUpResponse in catchUpReplies:    # Already decrypted

```

```

    if catchUpResponse.stateHash == correctHash:

```

```

        self.runningStateReplies.purge()    # Clear the response list

```

```

        send sign(<"runningStateRequest">) to catchUpResponse.replica

```

```

        timeout(self.replyTimeout):        # If we do not receive a response then try another

```

```

replica

```

```

        for stateResponse in self.runningStateReplies:

```

```

            # Decrypt using the replica's public key

```

```

            if decrypt(stateResponse) and if SHA256(stateResponse.runningState) ==

```

```

correctHash:    # Check if the hashes match

```

```

                correctRunningState = stateResponse.runningState

```

```

                break(2)          # Break 2 levels out of the outer loop

```

```

# Create a new configuration

```

```

initHist(self.replicaCount, self.config, hist = slot_op, runningState = correctRunningState)

```

```

}

```

**Replica {****fields:**

previousReplica  
nextReplica  
headReplica               *# the address to head Replica*  
history  
interReplicaKeys       *# public keys of all replicas*  
privateKey               *# the private key of this replica*  
status  
cache                    *# cache for storing the results*  
headFlag                *# boolean flag to check if replica is head or not*  
tailFlag                *# boolean flag to check if replica is tail or not*  
timeoutSpan             *# the timeout for retransmission requests*  
pendingRequests        *# requests whose result shuttle is pending*  
allowedOperations = ['query','update']  
maxSlotNum             *# the largest slot number ordered*  
runningState            *# contains the values of all objects in the replica*  
checkpointNum           *# the latest slot number which has been checkpointed*  
checkpointHash          *# the hash of the running state at the latest checkpoint*  
checkPointProof         *# latest checkpoint Proof*

**methods:**

checkpoint(checkPointSlot)  
  
handleAsyncRequest(msg, request, client)  
  
handleNewRequest(client, request)  
  
handleRetranmissionRequest(client, request)  
  
handleWedgeRequest(request)  
  
handleResultShuttle(result, resultProof)  
  
handleCheckpointShuttle(CheckPointProof)  
  
handleCheckpointReturnShuttle(CheckPointProof)  
  
handleCatchupRequest(request)  
  
handleRunningStateRequest(request)  
  
getProofsFromHistory(required\_slot)

```
orderCommand(order, orderProof, resultProof, sendFlag = True)
```

```
constructor(publicReplicaKeys, secretSigningKey, status, history, runningState):  
    self.status          = status  
    self.interReplicaKey  = publicReplicaKeys  
    self.privateKey       = secretSigningKey  
    self.maxSlotNum       = max(history) # Find the maximum slot number in the history  
    self.checkpointInterval = 100        # Denotes how often to checkpoint in terms of new  
slots ie. every 100 slots  
    self.runningState     = runningState
```

```
# This function performs a checkpoint up to the checkPointSlot
```

```
checkpoint(checkPointSlot):  
    # Update checkpointSlot number  
    self.checkpointNum = checkPointSlot
```

```
    # Hash the running state  
    self.checkPointHash = SHA256(self.runningState)    # calculate hash using SHA256
```

```
# Catch all message handler
```

```
handleAsyncRequest(msg, request, client):  
    switch msg:  
    'newRequest':  
        handleNewRequest(client, request)  
    'retransmission':  
        handleRetranmissionRequest(client, request)  
    'wedgeRequest':  
        handleWedgeRequest(request)  
    'catchUpRequest':  
        handleCatchupRequest(request)  
    'runningStateRequest':  
        handleRunningStateRequest(request)  
    'returnShuttle':  
        handleResultShuttle(result, resultProof)  
    'checkpointShuttle':  
        handleCheckpointShuttle(CheckPointProof)  
    'checkpointReturnShuttle':  
        handleCheckpointReturnShuttle(CheckPointProof)
```

```

        default:
            send sign(<"error", request.requestid>) to client    # Sign using replica's
private key

```

*# Asynchronous handler for normal order command requests*

*# Multi purpose handler to handle new requests*

**handleNewRequest**(client, request):

```

    if self.state != ACTIVE:                                # Check if we are active
        send sign(<"error", request.requestid>) to client    # Sign using replica's private key
    return

```

```

    if valid(client.clientKey) == False:                    # Check validity of client key
        send sign(<"error", request.requestid>) to client    # Sign using replica's private key
    return
    orderCommand(operations = request.operation, objectId = request.objectId)

```

*# Handler for retransmission requests*

**handleRetranmissionRequest**(client, request):

```

    if self.status == Immutable:
        send sign(<"error", request.requestid>) to client    # Sign using replica's private key
        return
    else if request.id in cache:                             # Sign using replica's private key
        send sign(<"retranmissionResponse", cache[request.id]>) to client
        return
    else:
        if self.headFlag == False:                          # forward the request to head
            sendRequestToReplica(headReplica, request)
            timeout(minutes = self.timeoutSpan):
                if request.id not in cache:                  # if no result shuttle obtained then reconfig
                    self.status = Immutable
                    send sign(<"reconfigure">) to Olympus    # Sign using replica's private key
                    return
                else
                    send sign(<"retranmissionResponse", cache[request.id]>) to client # Sign using
replica's private key
                    return
            else:
                # we are the head replica
                if request.id in cache:                      # if no result shuttle obtained then reconfig
                    send sign(<"retranmissionResponse", cache[request.id]>) to client # Sign using
replica's private key

```

```

        return
    else if request.id in pendingRequests:
        timeout(minutes = self.timeoutSpan):
            if request.id not in cache # if no result shuttle obtained then reconfig
                self.status = Immutable
                send sign(<"reconfigure">) to Olympus # Sign using replica's private key
                return
            else
                send sign(<"retranmissionResponse", cache[request.id]>) to client # Sign
using replica's private key
                return
        else:
            # similar to new request therefore calling handleRequest
            handleNewRequest(client, request)

```

# Handler for wedge requests

**handleWedgeRequest**(request):

```

    if (decrypt(request) == False)                # using Olympus public key
        return False                               # Do not respond

```

# Make sure our status is Immutable

self.status = Immutable

# Sign the data and send it to Olympus

send <"wedgeResponse", self.runningState, self.orderProofs, self.checkPointProof> to  
Olympus

# Handler for a result shuttle

**handleResultShuttle**(result, resultProof):

validFlag = True

for proof in resultProof:

if valid(proof) != True

validFlag = False

break

if validFlag == False:

# if the return Shuttle is not valid then ignore

return

else:

cache[result.id] = result # if it is valid then cache the result

pendingRequests.delete(result.id) # delete from the list of pending requests

# Handler to forward the checkpoint shuttle to the next replica

**handleCheckpointShuttle**(CheckPointProof):

```
# See proposeCheckpoint for checkpointTuple description
checkpointTuple = <self.checkpointNum, self.checkPointHash>
CheckPointProof.append(checkpointTuple)
if self.tailFlag:
    # Return to the previous replica
    handleCheckpointReturnShuttle(CheckPointProof)
else:
    # Forward to the next replica
    send <"checkpoint", CheckPointProof> to self.nextReplica
```

*# Handler to validate the return checkpoint shuttle*

**handleCheckpointReturnShuttle**(CheckPointProof):

```
# Validation
for checkpointTuple in CheckPointProof:
    if (checkpointTuple.checkpointNum != self.checkpointNum) or
(checkpointTuple.checkPointHash != self.checkPointHash):
        return Error                                # Do nothing

# Delete checkpointed proofs from history
checkpointSlot = CheckPointProof.checkpointTuple.checkpointNum
delete history[:checkpointSlot]                    # delete history till checkpoint

# store the returned checkPointProof from the tail
self.checkPointProof = CheckPointProof

# Return to the previous replica if you are not the head
if not self.headFlag:
    send <CheckPointProof> to self.previousReplica
```

*# Handler for catch up messages from Olympus*

**handleCatchupRequest**(request):

```
# Validate that this request came from Olympus. Decrypt using Olympus' public key
if not decrypt(request)
    return False
for order in request.catchupDifference:
    # Order command and do not forward to the next replica
    orderCommand(order, new OrderProof, new ResultProof, sendFlag = False)
```

```
stateHash = SHA-256(self.runningState)
```

key      send sign(<"catchupResponse", stateHash>) to Olympus      # Sign using replica's private

*# Handler for running state requests from Olympus*

**handleRunningStateRequest**(request):

*# Validate that this request came from Olympus. Decrypt using Olympus' public key*

if not decrypt(request)

    return False

send sign(<"runningStateResponse", self.runningState>) to Olympus

*# Gets all relevant proofs from the history*

*# start from current checkpoint and reach the required slot, using operations in history*

**getProofsFromHistory**(required\_slot):

    proof\_statements = self.history[self.checkpointNum:required\_slot]

    return proof\_statements

*# Relates to the orderCommand transition at the replica.*

*# Assumes we are active -> This is checked in the new request handler*

**orderCommand**(order, orderProof, resultProof, sendFlag = True):

*# Validate history*

if self.history has slot: # A slot for this command already exists. Cannot issue it twice

    return False

*# Generate a new slot if we are the head*

if self.headFlag == True:

    sNum = self.maxSlotNum+1

else

    sNum = order.slotNumber

*# Prevent holes in slot numbers*

if self.history does not contain sNum-1:checkpointNum: # We found a hole

    return False

*# Update maxSlotNum*

self.maxSlotNum = sNum

*# Apply running state*

result = runningState[request.objectId].apply(request.operation)

*# Create proofs*

```

        order = new OrderProof(slotNumber = sNum, operation = request.operation, proofs =
getProofsFromHistory(sNum))
        signedOrder = sign(key = self.privateKey, statement = order)
        orderProof.append(signedOrder)

        signedResult = sign(key = self.privateKey, statement = result)
        resultProof.append(signedResult)

# Initiate checkpoints after every X orders where X = self.checkpointInterval
        if maxSlotNum is a multiple of self.checkpointInterval:
            checkpoint(maxSlotNum)
            if self.headFlag # If we are the head replica, initiate the checkpoint shuttle
                handleCheckpointShuttle(new CheckPointProof())

        if not replica.tailFlag:
            send <"order", order, orderProof, resultProof> to self.nextReplica
        Else:
            # send result to client
            send <"result", requestID, result, resultProof> to request.client
            # send the return shuttle
            send <"ReturnShuttle", result, resultProof> to self.previousReplica
    }

```